

Implementation of Peer-to-Peer Energy Trading via Solana Smart Contracts in a Permissioned Environment

Chanthawat Kiriyaadee¹, and Suwannee Adsavakulchai^{1*}

¹Department of Computer Engineering and Artificial Intelligence, School of Engineering, University of the Thai Chamber of Commerce (UTCC), Bangkok, Thailand

Contact Email: 2410717302003@live4.utcc.ac.th, Suwannee_ads@utcc.ac.th

Thailand’s three state utilities — EGAT, MEA, and PEA — have pursued a National Energy Trading Platform (NETP) since 2017 for peer-to-peer (P2P) trading among rooftop-solar prosumers, but it remains in development, and the regulated buy-back affords prosumers no direct means of exchanging surplus within the local distribution grid. This paper evaluates a settlement layer for that role: six modular Anchor programs on a permissioned Solana runtime in which every high-frequency write is routed to a per-entity program-derived account (PDA), so transactions of distinct entities hold disjoint write sets and are scheduled in parallel by the Sealevel engine, ordered by Proof-of-History [1]. We measure compute, throughput and latency together, and find that they are not bound by the same thing. Compute is $O(1)$ and never binds: the steady telemetry write costs an identical 13,538 compute units (CU) at every scale of a 2,500× sweep from 80 to 200,000 meters and against 671,160 resident PDAs, while confirmation latency stays flat at p50 1.1–1.6 s (p95 ≤ 3.0 s) with zero delivery loss over 1.34 M first-attempt sends. Throughput is instead **associated** with lock footprint: ordering every measured path by the shared write-locked state it touches spans three orders of magnitude where the CU column does not — yet controlled experiments do not confirm the association as causal. Restoring the pre-fix lock footprint (the parent ZoneMarket re-declared writable, one account flag) leaves settlement throughput statistically unchanged from one fee payer to eight, and a previously reported ≈0.6 TPS proves to be a serial-await harness artifact. At ≈120 k CU a settle exhausts the block compute budget before its locks bind, so the write-set lever is not visible below a multi-node scheduler.

1. Introduction

Since 2017 Thailand’s state utilities have pursued a National Energy Trading Platform (NETP) for P2P trading among rooftop-solar prosumers, implemented on a Tendermint-based consortium blockchain [2]. This article presents a Solana-based architecture as an alternative execution layer, motivated by NETP’s own evaluation, in which Tendermint outperformed both Ethereum and Hyperledger Fabric under throughput testing: that result suggests further gains from a runtime whose execution is parallel rather than sequential. Such evaluations report throughput and compute cost as though computation were the scaling axis. The contribution here is to separate the axes: partition all hot-path state per entity, then measure compute, throughput and latency independently, so what binds each can be named rather than assumed.

2. The Per-Entity Partition

The settlement layer is six Anchor programs [3] (anchor-lang and anchor-spl 1.0.0) — **energy-token, governance, oracle, registry, trading, and treasury** — in Rust, compiled to SBF bytecode. Programs invoke one another through cross-program invocation (CPI), authority delegated to PDAs that sign without any private key. A PDA is the SHA-256 of its seed tuple, bump byte, program id and a domain-separation constant — accepted only if **not** a valid Ed25519 curve point, so no private key can exist and the owning program may sign as that address. A transaction declares in advance every account it touches and whether each is writable; the runtime locks per account, so two transactions with disjoint writable sets run on different cores and two writing one account serialise. Table 1 gives the hot-path partition under that rule.

Table 1: **Hot-path state partition.** Every high-frequency write targets its own program-derived account, so any two are Sealevel-parallelisable.

Hot path	Account	Seed tuple
meter telemetry	MeterState	[b"meter", meter_id] (32-byte id ceiling = the runtime’s seed limit)
order placement	Order	[b"order", authority, order_id]
settlement replay guard	OrderNullifier	[b>nullifier", user, order_id]
registration counters	RegistryShard × 16	[b"registry_shard", shard_id], shard = key[0] mod 16

Two consequences follow. First, global accounts — config, totals — are read-only on hot paths and **deliberately stale**; periodic admin instructions fold per-entity state back into them. Second, derivation is itself a budget item: `find_program_address` searches the bump byte downward from 255, a hash plus a curve-decompression check per candidate. The programs store the winning bump and thereafter validate with a single `create_program_address` — **1,651 CU against 12,136** on the aggregation path, a 7.3× saving on pure addressing.

3. Method

All experiments ran on one node (Apple M2, Agave 3.1.10) of a private solana-test-validator network, so they characterise SVM semantics — account locking and compute metering — not consensus or leader rotation [4]. Throughput is **client-observed confirmed goodput**: confirmed transactions per wall-clock second from burst start to last confirmation, covering the pipeline from client signing to confirmation visibility. Every non-confirmed transaction is attributed to one class — send-rejected (refused by the RPC node, no fee), guard-rejected (executed, failed a program check, fee paid and recorded on-chain), or expired — separating delivery loss from validation working as intended. Transactions are pre-signed against a cached blockhash and submitted raw without preflight or retry, then confirmed by bulk signature-status polling on a 1.5 s sweep under a 3,000-transaction in-flight window; latency is therefore poll-quantised (±1.5 s), an upper bound on the runtime’s own cost rather than a measurement of it. CU is machine-independent, read post-hoc from transaction metadata. The telemetry sweep submits one reading per meter from N distinct meters over two epochs at least 61 s apart — the rate-limit and monotonic-timestamp guards force the gap — separating one-off PDA creation from the recurring steady write. It covers $N = 80$ to 200,000 (1.34 M transactions) over two runs, the second unreset between scales (Table 3). The order-entry burst submits $N = 10^4$ orders over 16 round-robin zone shards.

4. Results and Discussion

Compute is $O(1)$ and never binds. Table 2 gives the instruction profile against the 200,000-CU default per-instruction budget; every control, telemetry and settlement-recording path costs at most ≈8% of it. Table 3 resolves the telemetry write by fleet size: with fixed-width meter identifiers the steady write is **13,538 CU in every cell of both runs** — 80 to 200,000 meters, a 2,500× range with zero drift — and the first write is likewise constant at 16,157 CU because it additionally initialises and rent-funds the PDA (one extra identifier byte costs ≈96 CU — the spread earlier variable-width measurements

showed as 13,337–13,634). Above the mode sits a deterministic ladder in exact 1,500-CU steps, geometrically decaying: find_program_address pays one 1,500-CU curve check per bump candidate and $\approx$ half of candidates fail, so half of all meters resolve at bump 255 and each further step halves — a per-meter constant, not noise. Settlement is the one expensive instruction at $\approx$ 121 k CU ($\approx$ 61%), dominated by native Ed25519 verification and instruction-sysvar introspection authorising the trade, not by settlement arithmetic.

Table 2: **Compute cost per instruction**, against the 200,000-CU default per-instruction budget. Machine-independent, read from transaction metadata; modes; PDA-deriving paths carry the bump ladder above the mode (see text).

Instruction	CU	% of 200 k
do_nothing (dispatch + signature-verify floor)	769	0.4%
treasury.record_settlement	3,302	1.7%
treasury.record_settlement_sharded	5,375	2.7%
trading.submit_limit_order_sharded	9,808	4.9%
oracle.submit_meter_reading (steady)	13,538	6.8%
oracle.submit_meter_reading (first, inits PDA)	16,157	8.1%
trading.settle_offchain_match	$\approx$ 121 k	$\approx$ 61%

Table 3: **Per-meter compute against fleet size**. Steady-state CU mode for oracle.submit_meter_reading, measured live, fixed-width meter ids. Run B continues run A's ledger without a reset, so its scales write against 161,160–471,160 already-resident PDAs and it ends holding 671,160. Every cell identical — the table is the $O(1)$ result.

Fleet size N (meters)	80	1,000	10,000	50,000	100,000	150,000	200,000
Steady write, run A (CU)	13,538	13,538	13,538	13,538	13,538	—	—
Steady write, run B (CU)	—	—	13,538	13,538	13,538	13,538	13,538

Compute is also $O(1)$ in accumulated state: run B's last 200,000 submissions cost the same against 471,160 resident accounts as the first scale's did against 161,160, and the ledger ends holding 671,160. The account store explains it: an update appends to an append-only file, and lookup goes through an in-memory pubkey $\rightarrow$ (slot, offset) index, so either write costs one probe and one append. What scales with account count — index memory, snapshot size, startup time — is a multi-node concern, invisible to per-transaction cost.

Throughput orders by lock footprint, but the mechanism does not confirm. Table 4 sorts every measured path by the shared writable state it touches, and that ordering — not the CU column — tracks the rate across three orders of magnitude. The causal test fails. Anchor's mut imposes a writability **requirement**, not a prohibition, so the shipped binary accepts the pre-fix lock footprint when the client declares ZoneMarket writable — one AccountMeta, verified on the wire. Over five repeats of 1,200 settles, arm order alternating, the writable arm returns 493 against 561 confirmed settles s^{-1} across eight fee payers and 447 against 446 across one. Even the shared fee-payer lock does not move throughput: conflicting transactions still land within one slot as successive entries, so a write lock costs wall-clock drain, not slot occupancy — and at $\approx$ 120 k CU per settle the block compute budget saturates before locks bind. We therefore withdraw the $2.7\text{--}3\times$ previously attributed to deleting this mut; both original figures are accounted for: a session's first burst runs $\approx$ 5–6 $\times$ slower than every later one (validator warm-up), and the $\approx$ 0.6 TPS once reported for single-payer settlement is a serial-await harness measuring its own round-trip (0.69 TPS, p50 463 ms reproduced). Order entry shares the caveat: the **cheaper** instruction returned 180 TPS against telemetry's 462 under a vestigial mut since dropped. Re-measured post-fix on the original 10^4 configuration, it lands all 10,000 orders at $\approx$ 1,000 TPS in every arm from one payer on one shard to 16 on 16 — a $5.7\times$ spanning the fix, a rewritten harness and a warm validator, so none of it attributes to the lock. Telemetry, holding only the gateway fee payer, sustains 234–296 TPS from 10,000 to 150,000 meters with no monotone trend, dipping to 182 in the final 200,000-meter epoch: not proportional, since one payer signs everything, and largely scale-independent, since per-meter PDAs never contend.

Table 4: **Throughput by lock footprint**. Paths ordered by the shared write-locked state they touch — an ordering, not a demonstrated cause (see text). Harnesses differ per row; rows are not mutually comparable. *pre-fix* = measured before the vestigial mut on zone_market was dropped.

Path	CU	TPS	Shared writable state
RPC read	—	$\approx$ 1,454	none — no consensus round-trip
meter ingest (10 k–200 k meters)	13,538	182–296	gateway fee payer
order entry (10 k orders, pre-fix)	9,884	180	fee payer + zone_market
order entry (10 k orders, re-measured)	9,808	915–1,028	fee payer + zone shard(s)
OLTP proxy (concurrency 5 $\rightarrow$ 40)	22–24 k	7.9 $\rightarrow$ 74.3	fee payer + market accounts
settlement, concurrent, 1–8 fee payers	$\approx$ 120 k	446–561	fee payer + energy_mint + collectors

Latency is set by block time, not by the partition. Confirmation holds p50 1.1–1.6 s and p95 $\leq$ 3.0 s across the whole 2,500 $\times$ fleet range — flat where a contended design would degrade — and order entry sits at p50 2,173 ms, p95 3,565 ms (measured pre-fix). Both are send $\rightarrow$ first-seen-confirmed upper bounds dominated by the $\approx$ 400–600 ms single-node block time, not measurements of runtime cost — visible in the read path, which needs no consensus round-trip and returns in under a millisecond. **Delivery loss is zero** across all 1.34 M first-attempt sends of the two sweeps — none refused at the RPC, none expired unconfirmed; the 0.47% residue is a deterministic anomaly-gate probe firing on the same meter indices every epoch, each landing on-chain as a fee-paid guard rejection.

5. Conclusion

With hot-path state partitioned per entity, neither computation nor per-entity contention bounds a P2P energy settlement layer: the steady write is $O(1)$ in fleet size and in resident account count, and latency is flat over a 2,500 $\times$ sweep. Lock footprint orders the measured paths more closely than instruction cost, but the ordering did not survive a controlled test as a mechanism: settlement throughput is unchanged from one fee payer to eight — at $\approx$ 120 k CU the settle exhausts the block compute budget before its locks bind — and the two strongest prior data points dissolve into validator warm-up and a serial-await harness. Shrinking declared write sets remains sound design; below a multi-node scheduler, a before/after on lock footprint chiefly measures its own harness. A single validator bounds execution and locking only; reproducing these results on a multi-validator consortium mapped to EGAT, MEA, and PEA would place both consensus and the scheduler under test, and suits UEC–ASEAN collaboration.

Acknowledgment

The authors thank the University of the Thai Chamber of Commerce (UTCC) for institutional support, and the organizers, the UEC ASEAN Research Center (UARC) and Multimedia University (MMU), for hosting this workshop.

References

- [1] A. Yakovenko, "Solana: A New Architecture for a High Performance Blockchain v0.8.13," 2018.
- [2] Metropolitan Electricity Authority (MEA) and Electricity Generating Authority of Thailand (EGAT) and Provincial Electricity Authority (PEA), "Development of the National Energy Trading Platform (NETP): Research and Development Project Report," technical report, Dec. 2019.
- [3] Coral / Anchor contributors, "Anchor: Solana's Sealevel runtime framework," [Online]. Available: <https://www.anchor-lang.com/>
- [4] Solana Foundation, "Tower BFT: Solana's High Performance Implementation of PBFT," Jul. 2019.